

Thesis Experiment Plan
OpenEM Based Video Filtering on Texas
Instruments TMS320C6678
Week 25

Risto Vuorio (84525R)
`risto.vuorio@aalto.fi`

June 16, 2015

Contents

1	Dynamic Workload Experiment on TMS320C6678	3
1.1	Outline	3
1.2	Goals	3
2	Texas Instruments TMS320C6678	4
2.1	Overview	4
2.2	Multicore Navigator	4
2.3	Development Environment	5
3	Filters	5
4	PREESM	5
5	OpenEM	6
5.1	Open Event Machine	6
5.2	Texas Instruments Implementation of OpenEM	6
6	Measurements	8
6.1	Desired Metrics	8
6.2	Implementation of the metrics	8
7	Construction	9
7.1	Video Filter Application Using PREESM	9
7.2	Video Filter Application Using OpenEM	10
7.3	Workload	10
8	Performance Simulation Environment	11
8.1	Performance Simulation Environment	11
8.2	Modeling the workload using PSE	11

1 Dynamic Workload Experiment on TMS320C6678

1.1 Outline

This document describes the construction of an OpenEM dynamic workload experiment. OpenEM implementation on Texas Instruments TMS320C6678 will be used as the runtime system. The workload will be a simple video filtering application consisting of couple of filters and dynamic number of streams.

The implementation steps of the experiment are represented in the following figure.

Iteration 1

- Filter application in PREESM
- Convert Filter application to OpenEM
- Implement tracing in both applications
- Model the OpenEM application with PSE
- Select a suitable workload for the applications
- Measure applications
- Compare measurements with the PSE estimates
- Check fitness of model to application behavior

Iteration 2

- Compare the model behavior to expected OpenEM behavior
- If the model behavior matches expected OpenEM behavior but not the real application behavior study the cause
 - Does the application utilize OpenEM according to specification?
 - Does OpenEM specification match the implementation on the platform
- If faults in either model or the application are found fix them

1.2 Goals

The objective of this experiment is to understand the behavior and performance of Texas Instruments implementation of Open Event Machine in processing a dynamic workload. The performance of an application implemented using OpenEM will be compared to the performance of a similar application implemented using a simpler multicore runtime. Open Event Machine utilizes a hardware accelerated packet communication system called The Multicore Navigator. The simpler runtime doesn't utilize hardware acceleration. The focus of this experiment is in understanding the effects of hardware accelerated scheduling vs. simple scheduling.

The OpenEM implementation of the workload application will be modeled using Performance Simulation Environment. The PSE simulation results are compared to the performance of the real world application. The simulation model created in PSE and validated on one real world application could be used to estimate the performance of similar applications on the hardware platform studied in this experiment or the performance of the application on similar hardware platforms.

2 Texas Instruments TMS320C6678

2.1 Overview

The hardware platform used in this experiment is Texas Instruments TMD-SEVM6678L TMS320C6678 Evaluation Module. The device belongs to the Keystone I family of multicore DSPs. The device has eight C6678 dsp cores. Each core can dispatch eight instructions every cycle. Communication between cores can be handled through hardware accelerated packet communication channel called Multicore Navigator or through shared memory using user defined locking scheme. For more details see the technical report by Hanhirona and Texas Instruments manuals.

Hanhirona, Jussi: TECHNICAL REPORT TI TMS320C6678 evaluation board characteristics, tools and analysis mechanisms for the purpose of static analysis
Texas Instruments: Multicore Fixed and Floating-Point Digital Signal Processor, tms320c6678.pdf

2.2 Multicore Navigator

The Multicore Navigator provides high speed packet communication on the device. The Multicore Navigator provides a hardware queue manager called the Queue Manager Subsystem (QMSS), packet DMA (PKTDMA) and multicore host notifications via interrupt generation.

In the Keystone I architecture there are two Packed Data Structure Processors (PDSP) in the QMSS. The PDSPs execute a firmware that can perform QMSS related functions. In this experiment firmware provided in the TI OpenEM implementation will be used. The firmware used handles the hardware queue management related to event scheduling.

The packets consist of descriptors and payload data. The packets reside in the memory and pointers to the packets are passed from core to core using the Multicore Navigator.

Texas Instruments: KeyStone Architecture Multicore Navigator, sprugr9h.pdf

2.3 Development Environment

Software is developed for the device using an IDE from Texas Instruments called Code Composer Studio. CCS version 5.2.1 is used. The Software Development Kit used for development for the device is called Multicore Software Development Kit (MCSDK). MCSDK version 2.1.2.5 is used.

CCS features a variety of debug, trace and analysis tools but to access performance counters in the hardware, libraries external to MCSDK are required. The needed trace libraries are distributed under collective name of CToolsLib. The newest versions of the libraries supporting the device are used.

CToolsLib download <https://gforge.ti.com/gf/project/ctoolslib/frs/?action=index>

3 Filters

A realistic workload application with capability for processing dynamic input is designed for the experiment. DSPs are commonly used for video stream processing. Processing video streams is a suitable workload for parallelization and dynamic number of inputs. Video filtering is a common and fairly simple task in video stream processing. For these reasons a video filtering application is selected as the workload for the experiment.

The workload consists of two common filters. The first is a sobel filter that is used in edge detection. The second is a gaussian filter that has a wide range of uses. Both filters are used in a canny edge detector which is a realistic workload for a multicore DSP. The real world users of the studied device would probably implement a complete edge detector rather than parts of it. However the workload in this experiment is deliberately kept simple to make it well analyzable and reasonably simple to implement.

The filters themselves are based on convolving the image data with a filter matrix and are thus straightforward to implement.

Gaussian Filter http://en.wikipedia.org/wiki/Gaussian_filter

Sobel Operator http://en.wikipedia.org/wiki/Sobel_operator

Canny Edge Detector http://en.wikipedia.org/wiki/Canny_edge_detector

4 PREESM

Actor Model programming is selected as the top-level programming model for the workload application. Actor Models are specifically designed for parallel programming and fit the Open Event Machine API layout well. Prototyping Actor Model programs is straightforward due to available graphical tools. Actor Model rapid prototyping tool PREESM will be utilized in the experiment

to build a simple version of the workload. PREESM was selected over comparable tools for the following reasons: 1. PREESM supports TMS320C6678 as a compilation target. 2. PREESM is under active development.

PREESM version 2.1.3 and Eclipse Luna Service Release 2 (4.4.2) will be used in the experiment.

Pelcat M. et. al. Preesm: A dataflow-based rapid prototyping framework for simplifying multicore DSP programming 10.1109/EDERC.2014.6924354

5 OpenEM

5.1 Open Event Machine

Open Event Machine is a runtime system for multicore platforms originally developed by NSN for the network dataplane. OpenEM is used to write dynamically load balanced applications with run-to-completion principle.

The key concepts of OpenEM are events, execution objects, queues and the scheduler. **Event** is the unit of communication in OpenEM. Events are typically used to carry data to be processed but can be data-less tokens as well. The event descriptors are allocated at the initialization of the Queues. **Execution objects** encapsulate the algorithm to execute when an event is received. **Queues** connect events (data) and execution objects (algorithms). Each queue is associated with one execution object and all queued events will be processed by this execution object. **Scheduler** moves allocated events to queues. **Dispatcher** is called by the user in a dispatch loop. Dispatcher calls the ‘receive’ function of the execution object of the connected queue.

5.2 Texas Instruments Implementation of OpenEM

The Texas Instruments implementation of Open Event Machine runtime system is included in the MCSDK (chapter 2.3). OpenEM version 1.0.0.2 is used in the experiment.

The OpenEM library is OS agnostic: the OpenEM dispatcher can be called from bare metal software or from an OS thread. The TI implementation of OpenEM Leverages the Multicore Navigator for hardware queues and packet communication. The scheduler always runs on a PDSP core and is defined by firmware specific to OpenEM.

The OpenEM execution model is described in figure 1. The event descriptors are allocated at OpenEM initialization in the free pool. The application allocates events from the free pool using *em_alloc*. Events can be allocated either at runtime or at initialization. After the allocated event has been populated it is sent to a queue using *em_send*. The queues are associated with specific execution objects at initialization and therefore they determine which EO will

process the events the queues. The scheduler running on a PDSP core (chapter 2.2) manages the queues (virtual queues mapped to QMSS hardware queues). The scheduling operations are carried out on an explicit request from a DSP core when a call to *ti_em_preschedule* is made. When one of the cores mapped to a specific execution object finishes executing its current event it makes a call to *ti_em_dispatch_once* which triggers a call to the *receive* function of the execution object. The execution object processes the data and makes a call to *em_free* to return the event to its free pool.

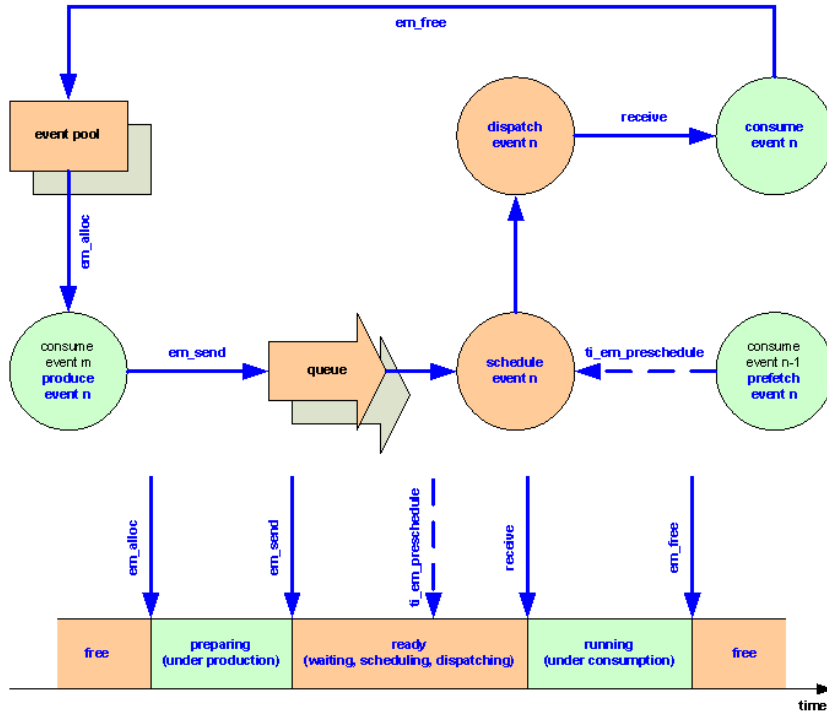


Figure 1: The OpenEM execution model

The concept of queue groups is used to specify which cores can execute which EOs. Queues are associated with EOs and each queue belongs to a queue group. Queue groups define which cores can execute the events from the queues in the group using a core mask.

The scheduler running on a PDSP core considers four scheduling criteria: **Priority**, **Atomicity**, **Locality** and **Order**. Each queue has a priority. The scheduler will always select events from queues with the highest priority. Queues are either atomic or parallel. Events from parallel queues can be scheduled on as many cores as are available for the queue. In contrast no event from an atomic

queue can be scheduled while another event from the same queue is executing. The scheduler tries to schedule events from certain queues on the same cores if possible. The last criterion for scheduling is order. If multiple events are eligible for scheduling, the event that has been in a queue longest will be scheduled.

Texas Instrument OpenEM White Paper, [ti.openem.white.paper.pdf](#)

Texas Instrument OpenEM User Guide, [ti.openem.user.guide.pdf](#)

Texas Instrument OpenEM Api Guide, [ti.openem.api.guide.pdf](#)

6 Measurements

6.1 Desired Metrics

Runtime performance measurements are carried out on both of the implemented applications. The high level metric which will tell us about the dynamic behavior of the OpenEM runtime model and the effect of the hardware acceleration is video frames per second. Finer grained metrics are implemented to understand the details of the Texas Instruments OpenEM implementation.

Area of interest for the finer grained metrics are the costs associated to the OpenEM runtime model. Even though the PREESM generated application is not an optimized application it will hopefully show the performance behavior to be expected from a simpler, non-hardware accelerated, “traditional” multicore implementation. Relating the PREESM application behavior to the OpenEM application behavior gives us a good view to the performance of the OpenEM runtime.

The specific time measurements to be carried out in addition to FPS are: time spent per actor, cost of switching from one actor to another on a core and time spent in the runtime. The time spent in the runtime can be further split to account for the time related to the scheduler running on a PDSP core and the time spent in the runtime on an actual C6678 core.

Other metrics will be implemented if the measurements described above will not yield accurate enough implementation about the OpenEM runtime behavior under dynamic loads.

6.2 Implementation of the metrics

The frames per second measurement is straightforward to implement by measuring the number of complete frames processed per time unit. The details of implementing the FPS measurement are not critical because completing a frame is relatively rare event in the execution of the application and calculating time spent per frame is simple.

Fine grained metrics will be implemented using Texas Instruments tools pro-

vided in CCS and CToolsLib. CCS provides a tool called System Analyzer which makes profiling the execution simple. System Analyzer can be used to easily measure time spent executing the actors. It is possible that the overhead caused by System Analyzer is too large to yield usable results. In that case the measurements will be carried out utilizing CToolsLib provided tools for capturing trace data and exporting the data through purpose built buffers.

On the first iteration all of the metrics are based on measuring execution time to keep things simple and fast to iterate. If some timing based measurements prove inaccurate or inefficient, they can be substituted with cycle-based measurements.

7 Construction

7.1 Video Filter Application Using PREESM

An actor network is constructed in PREESM that represents the video filter application. A network sketch is presented in Figure 2. To keep the model

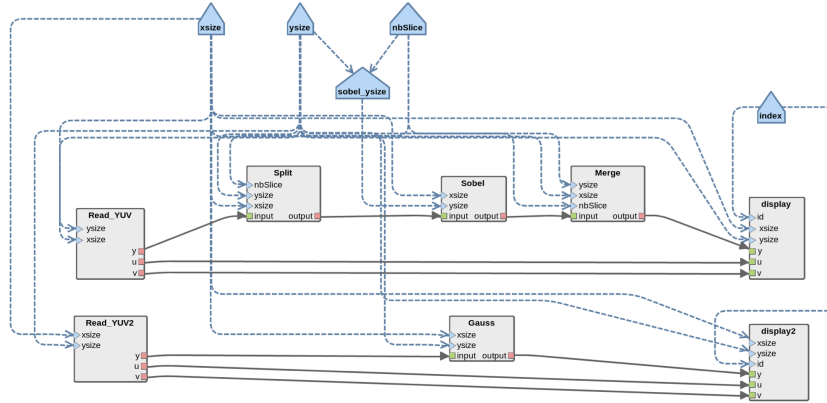


Figure 2: A Sketch of The Video Filter Actor Network

simple and the program well analyzable all of the processing paths in the network are independent. The shared dependencies in the sketch (marked with dashed lines) are constants related to the video stream and will change in the actual application.

Each filter implemented will have its own independent processing path.

The first actors on each of the processing paths load the video frames from memory and pass them to splitting actors. The splitting actor splits the frames to suitable number of splices to enable processing of the same video stream on multiple cores. The filter actor follows the splitting actors. Partial frames filtered in the filter actor are merged in to whole frames in the merge actors. The last actors on each processing path are sink actors that collect data from the execution and frees resources.

The filter actors will be run on all cores according to manual selection before execution. The cores on which the other actors will run on are to be determined when more knowledge about the execution is available.

7.2 Video Filter Application Using OpenEM

The video filter application first constructed as a PREESM actor model, will be implemented using OpenEM. The actors in the actor model map to execution objects in a straightforward manner. An execution object is allocated for each actor in the actor model. The events passed between the actors encapsulate the data in the video streams.

In contrast to the PREESM application where the schedule has to be changed for each number of video streams, the OpenEM implementation will not be changed for the different workloads.

The specifics of setting up the execution objects, number of queues attached to them, the priority of the queues and other specifics depend heavily on the final workload application and thus are not further explained here. The details are also subject to change based on the initial measurements to make the workload better reflect a realistic application.

7.3 Workload

The workload application needs to be kept simple to keep it analyzable. Complex parts of real applications unrelated to the parallel execution such as I/O are omitted from the workload implementation. Without I/O available the video streams need to be preloaded to the memory. CCS provides tools for preloading.

The purpose of the experiment is to understand the OpenEM behavior under dynamic load. Organizing the workload so that the number of streams could actually be dynamically changed at runtime will not be necessary if the cost of switching between streams can be estimated. A number of static loads will be constructed and measured on both of the applications.

8 Performance Simulation Environment

8.1 Performance Simulation Environment

Performance Simulation Environment (PSE) is a resource network based simulation environment that can be used to construct models of applications at different levels of abstraction. If a low abstraction level simulation model is desired, specific measurements of the runtime are needed. The measurements conducted on the second iteration will depend on the needs of the simulation model.

8.2 Modeling the workload using PSE

PSE will be used to construct a resource network model of the OpenEM based workload application. Building a resource network model will help us generalize the findings from the experiment.